
Guia 3 - Ejercicio 1

Table of Contents

Código fuente y resultados	1
Local Fourier analysis of sound	1
Exercise 1: (the solution is exo1.m)	7
Short time Fourier transform.	8
Audio Denoising	10
Exercise 2: (the solution is exo2.m)	13
Audio Block Thresholding	16

Código fuente y resultados

```
clear all;
close all;

numericaltour;
```

Local Fourier analysis of sound

First we load a sound, with a slight sub-sampling


```
n = 1024*16;
options.n = n;
[x,fs] = load_sound('bird', n);
```

You can actually play a sound.

```
sound(x(:)', fs);
```


We can display the sound.

```
figure();
clf;
plot(1:n,x);
axis('tight');
set_graphic_sizes([], 20);
title('Signal');
```



Local zooms on the sound show that it is highly oscillating.

```
figure();  
p = 512;  
t = 1:n;  
clf;  
sel = n/4 + (0:p-1);  
subplot(2,1,1);  
plot(t(sel),x(sel)); axis tight;  
sel = n/2 + (0:p-1);  
subplot(2,1,2);  
plot(t(sel),x(sel)); axis tight;
```




A good windowing function should balance both time localization and frequency localization.

```
t = linspace(-10,10,2048);
eta = 1e-5;
vmin = -2;
```


The block window has a sharp transition and thus a poor frequency localization.

```
h = double( abs(t)<1 );
hf = fftshift(abs(fft(h)));
hf = log10(eta+hf); hf = hf/max(hf);
clf;
subplot(2,1,1);
title('Block window');
plot(t, h); axis([-2 2, -.1, 1.1]);
subplot(2,1,2);
plot(t, hf); axis([-2 2, vmin, 1.1]);
title('Fourier transform');
```




A Hamming window is smoother.

```
h = cos(t*pi()/2) .* double(abs(t)<1);
hf = fftshift(abs(fft(h)));
hf = log10(eta+hf); hf = hf/max(hf);
clf;
subplot(2,1,1);
title('Hamming window');
plot(t, h); axis([-2 2, -.1, 1.1]);
subplot(2,1,2);
plot(t, hf); axis([-2 2, vmin, 1.1]);
title('Fourier transform');
```




A Hanning window has continuous derivatives.

```
h = (cos(t*pi()+1))/2 .* double(abs(t)<1);
hf = fftshift(abs(fft(h)));
hf = log10(eta+hf); hf = hf/max(hf);
clf;
subplot(2,1,1);
title('Hanning window');
plot(t, h); axis([-2 2, -.1, 1.1]);
subplot(2,1,2);
plot(t, hf); axis([-2 2, vmin, 1.1]);
title('Fourier transform');
```



A normalized Hanning window has a sharper transition. It has the advantage of generating a tight frame STFT, and is used in the following.

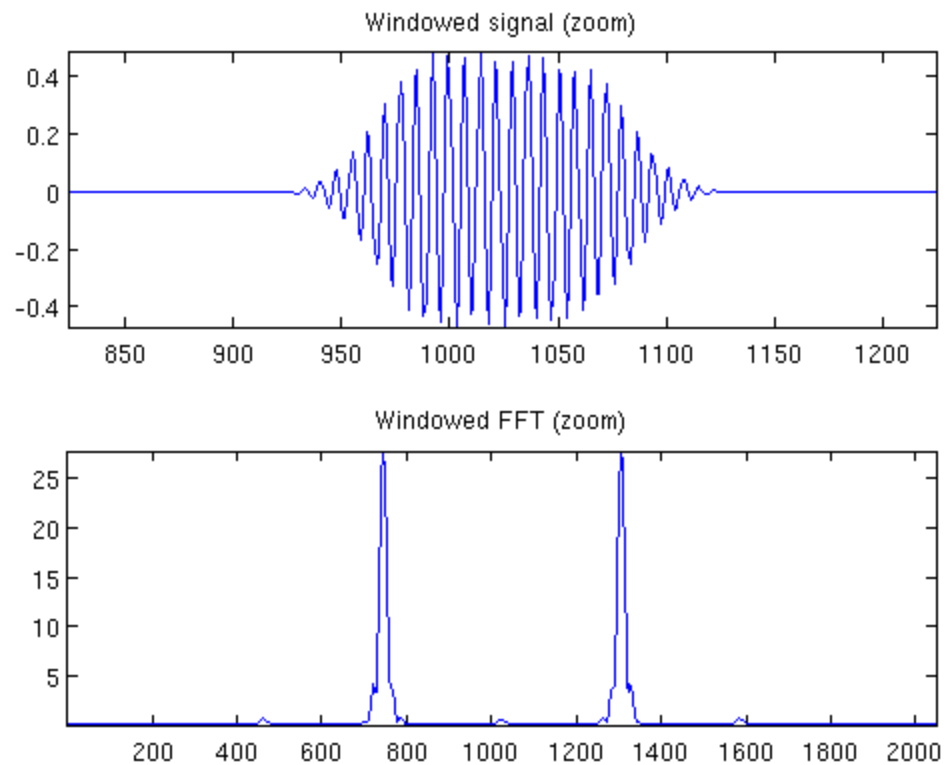
```
h = sqrt(2)/2 * (1+cos(t*pi())) ./ sqrt( 1+cos(t*pi()).^2 ) .* double(abs(t)<1);
hf = fftshift(abs(fft(h)));
hf = log10(eta+hf); hf = hf/max(hf);
clf;
subplot(2,1,1);
title('Normalized Hanning window');
plot(t, h); axis([-2 2, -.1, 1.1]);
subplot(2,1,2);
plot(t, hf); axis([-2 2, vmin, 1.1]);
title('Fourier transform');
```



Exercise 1: (the solution is exo1.m)

Compute the local Fourier transform around a point t_0 of x , which is the FFT (use the function `fft`) of the windowed signal $x.*h$ where h is smooth windowing function located around t_0 . For instance you can use for h a Gaussian bump centered at t_0 . To center the FFT for display, use `fftshift`.

```
figure();
t0 = 9500;
wl = length(hf);
wx = x(t0-floor(wl*0.5):t0+ceil(wl*0.5)-1).*h';
subplot(2,1,1),plot(wx);
axis([wl*0.5-200 wl*0.5+200 min(wx) max(wx)])
title('Windowed signal (zoom)')
subplot(2,1,2),plot(fftshift(abs(fft(wx))));
title('Windowed FFT (zoom)')
axis('tight')
```



Short time Fourier transform.

The only parameters of the transform are the size of the window and the overlap.

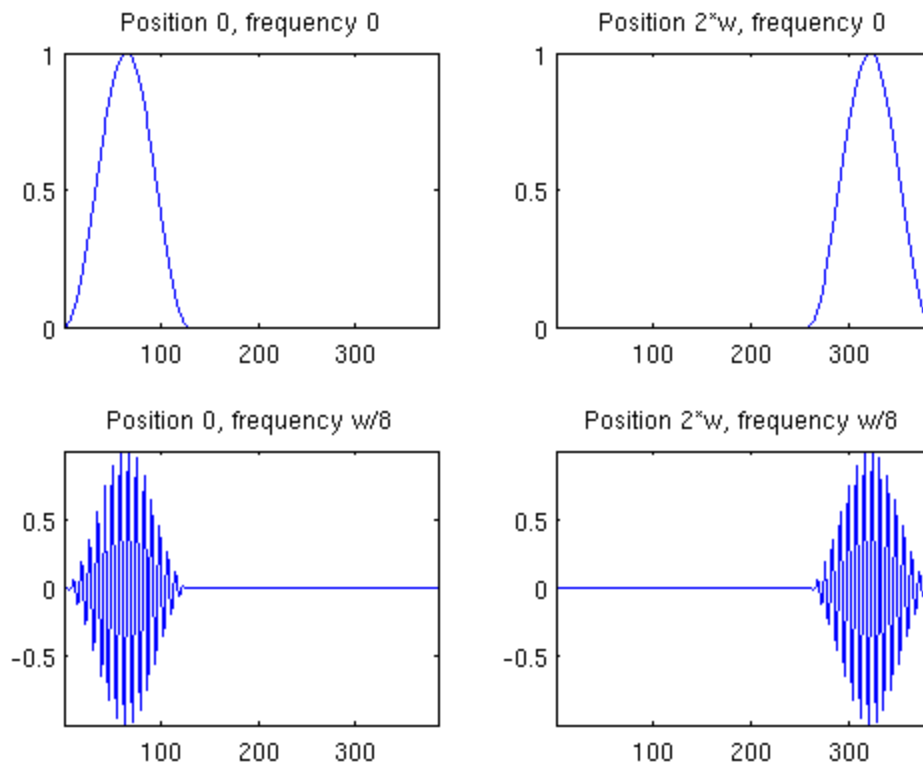
```
% size of the window
w = 64*2;
% overlap of the window
q = w/2;
```

Gabor atoms are computed using a Hanning window. The atoms are obtained by translating in time and in frequency (modulation) the window.

```
t = 0:3*w-1;
t1 = t-2*w;
f = w/8;
% Position 0, frequency 0.
g1 = sin( pi*t/w ).^2 .* double(t<w);
% Position 2*w, frequency 0.
g2 = sin( pi*t1/w ).^2 .* double( t1<w & t1>=0 );
% Position 0, frequency w/8
g3 = g1 .* sin( t * 2*pi/w * f);
% Position 2*w, frequency w/8
g4 = g2 .* sin( t * 2*pi/w * f);
% display
figure()
clf;
```



```
subplot(2,2,1);
plot(g1); axis('tight');
title('Position 0, frequency 0');
subplot(2,2,2);
plot(g2); axis('tight');
title('Position 2*w, frequency 0');
subplot(2,2,3);
plot(g3); axis('tight');
title('Position 0, frequency w/8');
subplot(2,2,4);
plot(g4); axis('tight');
title('Position 2*w, frequency w/8');
```




We can compute a spectrogram of the sound to see its local Fourier content. The number of windowed used is $(n-\text{noverlap})/(\text{w}-\text{noverlap})$

```
S = perform_stft(x,w,q, options);
```

To see more clearly the evolution of the harmonics, we can display the spectrogram in log coordinates. The top of the spectrogram corresponds to low frequencies.

```
% display the spectrogram
clf; imageplot(abs(S)); axis('on');
% display log spectrogram
plot_spectrogram(S,x);
```



The STFT transform is decomposing the signal in a redundant tight frame. This can be checked by measuring the energy conservation.

```
% energy of the signal
e = norm(x, 'fro').^2;
% energy of the coefficients
eS = norm(abs(S), 'fro').^2;
disp(strcat(['Energy conservation (should be 1)=', num2str(e/eS)]));
```

Energy conservation (should be 1)=1

One can also check that the inverse transform (which is just the transposed operator - it implements exactly the pseudo inverse) is working fine.

```
% one must give the signal size for the reconstruction
x1 = perform_stft(S,w,q, options);
disp(strcat(['Reconstruction error (should be 0)=', ...
    num2str( norm(x-x1, 'fro')./norm(x, 'fro') ) ]));
```

Reconstruction error (should be 0)=2.2398e-16

Audio Denoising

One can perform denoising by a non-linear thresholding over the transformed Fourier domain.

First we create a noisy signal


```
sigma = .2;
snoise = randn(size(x))*sigma;
xn = x + snoise;
```

Play the noisy sound.

```
sound(xn,fs);
```


Display the Sounds.

```
clf;
figure()
subplot(2,1,1);
plot(x); axis([1 n -1.2 1.2]);
set_graphic_sizes([], 20);
title('Original signal');
subplot(2,1,2);
plot(xn); axis([1 n -1.2 1.2]);
set_graphic_sizes([], 20);
title('Noisy signal');
```



One can threshold the spectrogram.

```
% perform thresholding
Sn = perform_stft(xn,w,q, options);
SnT = perform_thresholding(Sn, 2*sigma, 'hard');
% display the results
figure()
subplot(2,1,1);
plot_spectrogram(Sn);
subplot(2,1,2);
plot_spectrogram(SnT);
```



Exercise 2: (the solution is exo2.m)

A denoising is performed by hard or soft thresholding the STFT of the noisy signal. Compute the denoising SNR with both soft and hard thresholding, and compute the threshold that minimize the SNR. Remember that a soft thresholding should be approximately twice smaller than a hard thresholding. Check the result by listening. What can you conclude about the quality of the denoised signal ?

```
x2h = perform_stft(SnT,w,q, options);

SnT = perform_thresholding(Sn, sigma, 'soft');
figure()
subplot(2,1,1);
plot_spectrogram(Sn);
subplot(2,1,2);
plot_spectrogram(SnT);

% Reconstruyo

x2s = perform_stft(SnT,w,q, options);
sound(x2h,fs);
sound(x2s,fs);

close all;
Tsigma = 0.8:0.05:2.5;
snr_soft = 0.*Tsigma;
```

```

snr_hard = 0.*Tsigma;

% Acá está más explicado:
% https://www.ceremade.dauphine.fr/~peyre/numerical-tour/tours/denoisingwav\_2\_wave
% universal threshold rule
%  $T = \sigma \sqrt{2 \log(N)}$ 
for k=1:1:length(Tsigma)

    SnTh = perform_thresholding(Sn, Tsigma(k)*sigma, 'hard');
    x2h = perform_stft(SnTh,w,q, options);
    snr_hard(k) = snr(x,x2h);

    % Sugerencia del autor del tutorial: Remember that a soft thresholding
    % should be approximately twice smaller than a hard thresholding.
    SnTs = perform_thresholding(Sn, Tsigma(k)*sigma*0.5, 'soft');
    x2s = perform_stft(SnTs,w,q, options);
    snr_soft(k) = snr(x,x2s);

end

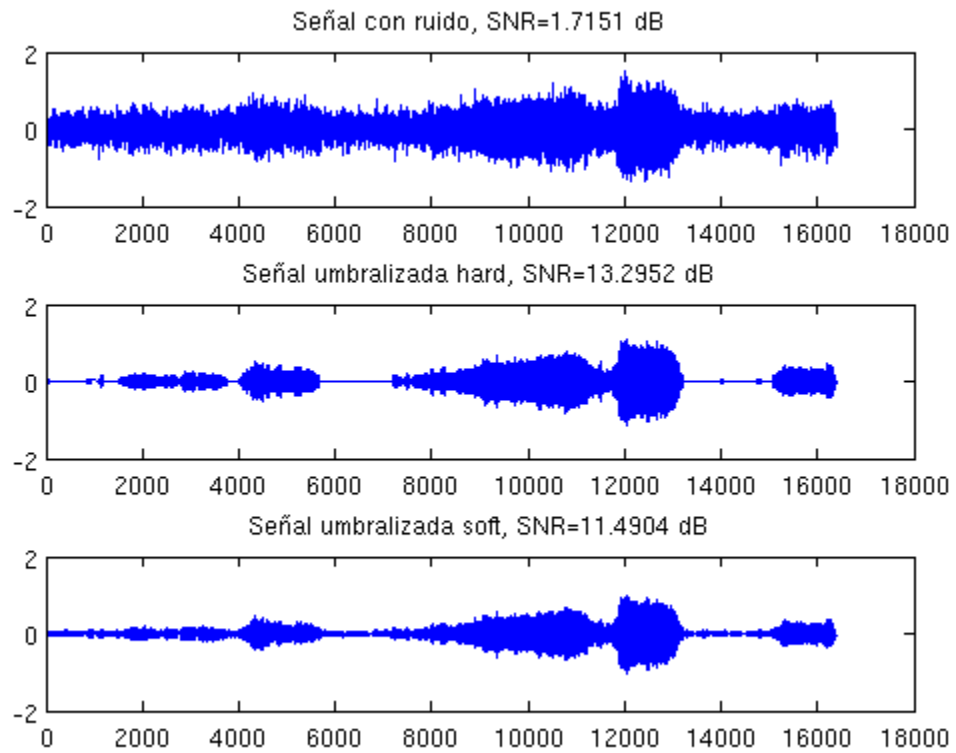

figure()
clf;
plot(Tsigma,snr_soft,Tsigma,snr_hard);
legend('Soft','Hard');

% Veo los mejores umbrales
[ms,poss] = max(snr_soft);
[mh,posh] = max(snr_hard);

SnTh = perform_thresholding(Sn, Tsigma(posh)*sigma, 'hard');
x2h = perform_stft(SnTh,w,q, options);
SnTs = perform_thresholding(Sn, Tsigma(poss)*sigma*0.5, 'soft');
x2s = perform_stft(SnTs,w,q, options);

figure()
subplot(3,1,1),plot(xn);
title(sprintf('Señal con ruido, SNR=%.4f dB',snr(x,xn)));
subplot(3,1,2),plot(x2h);
title(sprintf('Señal umbralizada hard, SNR=%.4f dB',snr(x,x2h)));
subplot(3,1,3),plot(x2s);
title(sprintf('Señal umbralizada soft, SNR=%.4f dB',snr(x,x2s)));

```



Audio Block Thresholding

It is possible to remove musical noise by thresholding blocks of STFT coefficients.

Denoising is performed by block soft thresholding.

```
% perform thresholding
Sn = perform_stft(xn,w,q, options);
SnT = perform_thresholding(Sn, sigma, 'block');
% display the results
subplot(2,1,1);
plot_spectrogram(Sn);
subplot(2,1,2);
plot_spectrogram(SnT);

% Exercice 4: (the solution is exo4.m) Try for various block sizes and
% report the best results.

snr_block = zeros(4,4);
for k=1:1:4
    for l=1:1:4
        options.block_size = [k, l];
        SnTb = perform_thresholding(Sn, sigma, 'block', options);
        x2b = perform_stft(SnTb,w,q, options);
        snr_block(k,l) = snr(x,x2b);
    end
end

figure();
imagesc(snr_block)
colormap(gray)




[mb, posb] = max(snr_block(:));
[k,l] = ind2sub(size(snr_block),posb);


options.block_size = [k, l];

SnTb = perform_thresholding(Sn, sigma, 'block', options);
x2b = perform_stft(SnTb,w,q, options);


% display the results
figure()
subplot(2,1,1);
plot_spectrogram(Sn);
subplot(2,1,2);
plot_spectrogram(SnTb);
title('Denoise process with optimal block size');

figure()
subplot(2,1,1),plot(xn);
title(sprintf('Señal con ruido, SNR=%.4f dB',snr(x,xn)));
subplot(2,1,2),plot(x2b);
title(sprintf('Señal umbralizada hard, SNR=%.4f dB',snr(x,x2b)));
```








Denoise process with optimal block size



Señal con ruido, SNR=1.7151 dB



Señal umbralizada hard, SNR=14.1052 dB



Published with MATLAB® R2013a